

🎓 Lecture 07 – What is Vector Database | Internal Implementation of Vector DataBase

(Expert-level, engaging, and complete — by Harshal Chauhan 🌟)

🧩 1 Introduction – Kyun Chahiye Vector Databases?

Traditional databases (SQL/NoSQL) sirf **exact match** search karte hain, jabki **Vector Databases** “meaning-based” search karte hain 🔍

⚙️ Comparison:

📦 **SQL/NoSQL**: “Rohit Negi 9 ke comments lao” → Exact Match ✅

🧠 **Vector DB**: “Rohit Negi 9 jaisi profiles lao” → Semantic Similarity ✅

💬 *Example*: “What is an array?” ≈ “Array kya hota hai?” → Same Results 🔄

🛒 2 Real-World Example – Blinkit Recommendation

User adds “**Onion**” to cart → Vector DB recommends **similar items**

✅ Recommended (Similar Vectors)

- Tomato → [0.8, 0.2, 0.0, 0.8, 0.2]
- Dhaniya → [0.7, 0.1, 0.0, 0.9, 0.0]
- Nimbu → [0.1, 0.8, 0.1, 0.7, 0.3]

❌ Not Recommended (Different Vectors)

- Banana → [0.1, 0.9, 0.1, 0.1, 0.9]
- Protein → [0.0, 0.1, 0.9, 0.1, 0.7]
- Creatine → [0.0, 0.0, 0.95, 0.0, 0.1]

💡 *System meaning ke basis par recommend karta hai, na ki sirf text match pe!*

**3**

Vector Similarity – Core Concept

◆ Formula (Euclidean Distance):

$$\text{Distance} = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

🧠 Example:

Onion [2,3] & Tomato [3,4]

$$\rightarrow \sqrt{(3-2)^2 + (4-3)^2} = \sqrt{2} \approx 1.4$$

📌 Rule:

👉 *Jitni chhoti distance, utni zyada similarity!*

**4**

Brute Force Approach – The Slow Killer ❌

🔍 Query:

Find 10 nearest neighbors of “Onion”

Steps:

1. Onion vector = [0.9, 0.1, 0.0, 0.9, 0.1]
2. Compare with *1 billion* vectors
3. Calculate distance for each
4. Sort results
5. Return top 10

⚠️ Result:

- ✅ Accuracy: 100%
- 🐢 Speed: Extremely slow
- ❤️ Bad User Experience

💡 *Brute Force = Perfect but Painfully Slow*

**5**

ANN – Approximate Nearest Neighbor

🎯 Smart Trade-off

Accuracy**Speed**

100% →
90–95%

🚀 100x
Faster

Example:

10 recommendations → 9 correct + 1 slightly off = still worth it!

💡 *Speed > Perfection in real-world systems.*

🌟 6 ANN Algorithms – Deep Dive

◆ 6.1 IVF – Inverted File Index

Phase 1: Indexing

- Step 1: **Find Centroids** using *K-Means*
- Step 2: Assign each vector to nearest centroid (cluster)

Phase 2: Searching


- Query “Grapes” → Compare with 100 centroids
- Choose nearest (say C2) → Search only in that cluster

➡ 100x faster!

💡 Multi-Centroid Strategy:

Query top 3 closest centroids → better accuracy 🎯

📊 Text Figure:



Each C = centroid, connected to its vectors.

🌳 6.2 KD-Tree – Binary Space Partitioning

Process:

- Divide data recursively using X & Y splits
- Create a decision tree
- Query travels through branches based on coordinates

Example:

Query Q(13,8)

→ Root split at x=11 → Go Right

→ y=10 → Go Left

→ Compare only with nearby nodes

🔥 Efficient for low-dimensional data.

Text Figure:

```

      (x=11)
    /      \
  Left      Right
(y=10)      (y=15)
 /  \      /  \
[Pts] [Pts] [Pts] [Pts]
```

6. HNSW – Hierarchical Navigable Small Worlds


Concept: Multi-layer graph structure


Layers:

- Layer 0: All nodes connected to 3 nearest
- Layer 1: Random nodes promoted
- Layer 2: Even fewer nodes
- Top Layer: Only 2–3 nodes

Search Process:

1. Start at top layer
2. Find nearest
3. Move layer by layer downward
4. Continue till Layer 0

 **Alpha=3** → Check 3 neighbors per level

 *Fastest and most accurate ANN algorithm.*

Text Figure:

```

Layer 2:   A – B
           |
Layer 1:   A–C–B–D
           |
Layer 0:   A–E–C–F–B–G–D
```

6.4 PQ – Product Quantization

Problem:

1 Billion vectors = 6.1 TB 🤯

Solution (Compression):

- Step 1: Split vector into 12 chunks
- Step 2: Create 256 centroids (codebook)
- Step 3: Replace chunks with centroid indices → only 12 bytes!

🔧 Memory Saved:

6.1 TB → 12 GB (≈500x reduction!)

PQ Search:

- Split query into chunks
- Compare with precomputed distance tables

⚡ *Super-fast lookup!*

Text Figure:

Vector [v1 v2 v3 v4 ... v12]

↓ Split into chunks

[] [] [] ... []

↓ Quantized (Centroid IDs)

[05][19][02]...[77]

6.5 IVF + PQ – Hybrid Approach

Combo of:

1. **IVF**: Smart clustering
2. **PQ**: Compression within each cluster

🎯 Result:

- High speed
- Low memory
- Great accuracy

💡 *Used widely in billion-scale systems.*

🎯 Text Figure:

[Clusters]

- C1 → PQ compressed vectors
- C2 → PQ compressed vectors
- C3 → PQ compressed vectors

📊 7 Performance Comparison

Algorithm m	Spe ed 🚀	Accura cy 🎯	Memo ry 📁	Ideal Use
🏆 HNSW	9/10	10/10	8/10	Maximum Performance
⚡ IVF + PQ	8/10	8/10	4/10	Large Datasets
⚖️ IVF Only	6/10	6/10	5/10	Balanced Needs
🌳 KD-Tree	4/10	5/10	6/10	Low Dimensions
📊 Brute Force	1/10	10/10	9/10	Small Datasets

📁 8 Vector Database Ecosystem

🧠 Native Vector Databases

Database	Algorithm	Feature
Milvus	HNSW, IVFPQ	Open-source, scalable
Pinecone	HNSW	Managed, high performance
Weaviate	HNSW	GraphQL + Knowledge Graph
Qdrant	HNSW + Quantization	Built in Rust, fast

Integrated Solutions

- PostgreSQL → **pgvector extension**
- Elasticsearch → **Text + Vector search**
- Redis → **In-memory + Vector support**

Foundation Library

- **Faiss (Meta)** → HNSW, IVFPQ (Industry Standard)

9 Vector Database Storage Structure

ID → Unique key

Vector → Similarity search

Metadata → Filters (category, price, etc.)

Example:

```
{ id: "product_456", vector: [0.98, 0.23, ...],  
  metadata: {"name": "Red Shoes", "price": 89.99} }
```

10 Search Types

1. **Vector Similarity Search** → “Onion jaisi items lao”
2. **Exact ID Search** → “Product_456 lao”
3. **Hybrid Search** → “Onion jaisi items under ₹50 lao”

💡 *Hybrid = Vector + Metadata combo*






11 Algorithm Selection Guide

Case → **Recommended**

- ✅ Best Performance → HNSW
- 💾 Low Memory → IVF + PQ
- ⚖️ Balanced → IVF Only
- 🌲 Low Dimension → KD-Tree

1 Ultimate Summary

Key Insights

1.  Brute Force → Slow
2.  ANN → Fast + Smart
3.  HNSW → Best performance
4.  IVF+PQ → Memory efficient
5.  Hybrid → Practical choice

 *Vector DBs samajhte hain meaning, sirf words nahi!*

1 Implementation Roadmap

Type → **Tool**

Startup → Pinecone


Enterprise → Milvus


SQL Users → PostgreSQL + pgvector

Developers/Researchers → Faiss

1 Final Takeaway

- ◆ **Traditional DBs** → Exact Match
- ◆ **Vector DBs** → Semantic Understanding
- ◆ **ANN Algorithms** → Fast + Smart
- ◆ **Hybrid Systems** → Best Real-World Performance

 *In short:*

“Vector DBs think like humans — they understand meaning, not just text.” 

 **Made by Harshal Chauhan | Expert Notes for Visual Learning** 